

Importing The Important Modules From The Library

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
```

Load The Dataset

```
df=pd.read_csv("Titanic-Dataset.csv")#Load the dataset
```

```
df.head()#Check the head
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	SibSp	\	Name	Sex	Age
0			Braund, Mr. Owen Harris	male	22.0
1					
1	1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1					
2			Heikkinen, Miss. Laina	female	26.0
0					
3			Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
1					
4			Allen, Mr. William Henry	male	35.0
0					

	Parch		Ticket	Fare	Cabin	Embarked
0	0		A/5 21171	7.2500	NaN	S
1	0		PC 17599	71.2833	C85	C
2	0	STON/O2.	3101282	7.9250	NaN	S
3	0		113803	53.1000	C123	S
4	0		373450	8.0500	NaN	S

```
# check the total row and the column of the dataset
```

```
df.shape
```

```
(891, 12)
```

Data Cleaning

```
#check the missing value
```

```
mv=df.isnull().sum()
```

```
print(mv)
```

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

```
# Fill missing values in 'Age' with their respective means
```

```
df['Age'].fillna(df['Age'].mean(), inplace=True)
```

```
# Let's also fill 'Embarked' with its mode, as it's categorical
```

```
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

```
# Drop 'Cabin' due to too many missing values
```

```
df.drop(columns=['Cabin'], inplace=True)
```

```
# Check if any missing values remain
```

```
ms = df.isnull().sum()
```

```
ms
```

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age             0
SibSp           0
Parch           0
Ticket          0
Fare            0
Embarked        0
dtype: int64
```

```
df.head()
```

```
   PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
```

3	4	1	1
4	5	0	3

SibSp \	Name	Sex	Age
0	Braund, Mr. Owen Harris	male	22.0
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1	Heikkinen, Miss. Laina	female	26.0
2	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
3	Allen, Mr. William Henry	male	35.0

	Parch	Ticket	Fare	Embarked
0	0	A/5 21171	7.2500	S
1	0	PC 17599	71.2833	C
2	0	STON/O2. 3101282	7.9250	S
3	0	113803	53.1000	S
4	0	373450	8.0500	S

```
df.info() #for knowing the information
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age            891 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Embarked        891 non-null   object
dtypes: float64(2), int64(5), object(4)
memory usage: 76.7+ KB
```

Encode the dataset

```
print(df.columns)
```

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age',
      'SibSp',
```

```
'Parch', 'Ticket', 'Fare', 'Embarked'],
dtype='object')
```

```
#for binary categories like Sex put Male equals to 0 and female equals to 1
```

```
df['Sex'] = df['Sex'].map({'male': 0, 'female': 1})
```

```
# for Embarked, which has more than two values like C, Q, S so we use dummy variable
```

```
df = pd.get_dummies(df, columns=['Embarked'], drop_first=True)
```

```
#It automatically dropped the first category (likely 'C') to avoid the dummy variable trap (multicollinearity).
```

```
# Embarked_Q = 1 → Passenger embarked at Queenstown
```

```
# Embarked_S = 1 → Passenger embarked at Southampton
```

```
# If both are 0, it means the passenger embarked at Cherbourg (C).
```

```
df.head()
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp
Parch \				
0	Braund, Mr. Owen Harris	0	22.0	1
0				
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	38.0	1
0				
2	Heikkinen, Miss. Laina	1	26.0	0
0				
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	35.0	1
0				
4	Allen, Mr. William Henry	0	35.0	0
0				

	Ticket	Fare	Embarked_Q	Embarked_S
0	A/5 21171	7.2500	0	1
1	PC 17599	71.2833	0	0
2	STON/O2. 3101282	7.9250	0	1
3	113803	53.1000	0	1
4	373450	8.0500	0	1

Find the Correlation

```
# Correlation with Survived
```

```
correlation = df[['Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Sex',
```

```
'Embarked_Q', 'Embarked_S', 'Survived']].corr()
print(correlation['Survived'].sort_values(ascending=False)) #
ascending = false means order will be highest to lowest
```

```
Survived      1.000000
Sex           0.543351
Fare          0.257307
Parch         0.081629
Embarked_Q    0.003650
SibSp        -0.035322
Age          -0.069809
Embarked_S   -0.149683
Pclass       -0.338481
Name: Survived, dtype: float64
```

Split the dataset

```
# Define features (X) and target (y)
X = df[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare',
'Embarked_Q', 'Embarked_S']]
y = df['Survived']

#Split data (70% train, 30% test)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# check the shape of the training data
X_train.shape

(623, 8)

# check the shape of the testing data
X_test.shape

(268, 8)
```

Model Buiding

At the first time I use logistic regression model and then I use the random forest model for better model prediction making

Logistic Regression Model

```
# Initialize and train the model
model = LogisticRegression(max_iter=1000, random_state=42)
model.fit(X_train, y_train)

# Predict on test set
y_pred = model.predict(X_test)
```

Test the model

```
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
```

```
# Evaluate Logistic Regression
```

```
print("Logistic Regression Performance:")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test,
y_pred))
```

```
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Logistic Regression Performance:

Accuracy: 0.8097014925373134

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.87	0.84	157
1	0.79	0.73	0.76	111
accuracy			0.81	268
macro avg	0.81	0.80	0.80	268
weighted avg	0.81	0.81	0.81	268

Confusion Matrix:

```
[[136  21]
 [ 30  81]]
```

Checking the model with new data

```
new_passenger = pd.DataFrame({
    'Pclass': [3], 'Sex': [0], 'Age': [25], 'SibSp': [0], 'Parch':
[0],
    'Fare': [7.5], 'Embarked_Q': [0], 'Embarked_S': [1]
})
```

```
prediction = model.predict(new_passenger)
```

```
print("Survived" if prediction[0] == 1 else "Did Not Survive")
```

Did Not Survive

```
new_passenger = pd.DataFrame({
    'Pclass': [3], 'Sex': [1], 'Age': [15], 'SibSp': [0], 'Parch':
[0],
    'Fare': [8.5], 'Embarked_Q': [1], 'Embarked_S': [1]
})
```

```
prediction = model.predict(new_passenger)
```

```
print("Survived" if prediction[0] == 1 else "Did Not Survive")
```

Survived

Random Forest

```
from sklearn.ensemble import RandomForestClassifier

# Initialize and train Random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Predict on test set
y_pred_rf = rf_model.predict(X_test)

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
# Optionally, evaluate Random Forest
print("\nRandom Forest Performance:")
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
print("\nClassification Report:\n", classification_report(y_test,
y_pred_rf))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
```

Random Forest Performance:
Accuracy: 0.7910447761194029

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.84	0.83	157
1	0.76	0.72	0.74	111
accuracy			0.79	268
macro avg	0.79	0.78	0.78	268
weighted avg	0.79	0.79	0.79	268

Confusion Matrix:

```
[[132  25]
 [ 31  80]]
```

```
new_passenger = pd.DataFrame({
    'Pclass': [3], 'Sex': [0], 'Age': [25], 'SibSp': [0], 'Parch':
[0],
    'Fare': [7.5], 'Embarked_Q': [0], 'Embarked_S': [1]
})
prediction = model.predict(new_passenger)
print("Survived" if prediction[0] == 1 else "Did Not Survive")
```

Did Not Survive

```
new_passenger = pd.DataFrame({
    'Pclass': [3], 'Sex': [1], 'Age': [15], 'SibSp': [0], 'Parch':
    [0],
    'Fare': [8.5], 'Embarked_Q': [1], 'Embarked_S': [1]
})
prediction = model.predict(new_passenger)
print("Survived" if prediction[0] == 1 else "Did Not Survive")
Survived
```

Conclusion

When to choose Logistic Regression (accuracy = 0.81):

Simpler, faster, and easier to interpret.

Works well if the relationship between features and outcome is linear.

Less prone to overfitting with smaller datasets.

Good if you want to explain results clearly.

When to choose Random Forest (accuracy = 0.79):

Handles non-linear relationships and interactions better.

More robust to outliers and noisy data.

Better if your data is large, complex, or includes categorical features.

Can rank feature importance.

So, I think logistic regression is slightly better than random forest model. I always prefer the logistic regression model